Where the Moats Are

SmallPE Research Note – AI Competitive Advantage, Late 2026

SmallPE Investment Team

October 26, 2026

Thesis

The AI capability race split into two separate races, and conflating them is the most common diligence error we see today. **Open-weight models won reach.** Led by Chinese labs – Qwen, DeepSeek, Kimi, GLM, Ling – after Meta's 2026 retreat to closed development, open models now carry the large majority of commodity inference volume at a fraction of closed cost. **Closed frontier labs won revenue.** OpenAI, Anthropic and Google tripled enterprise spend capture in a year by monetizing a narrow, persistent capability lead on the hard work. We think the second race determines where investable value accrues – and the asset worth underwriting is no longer the model at all.

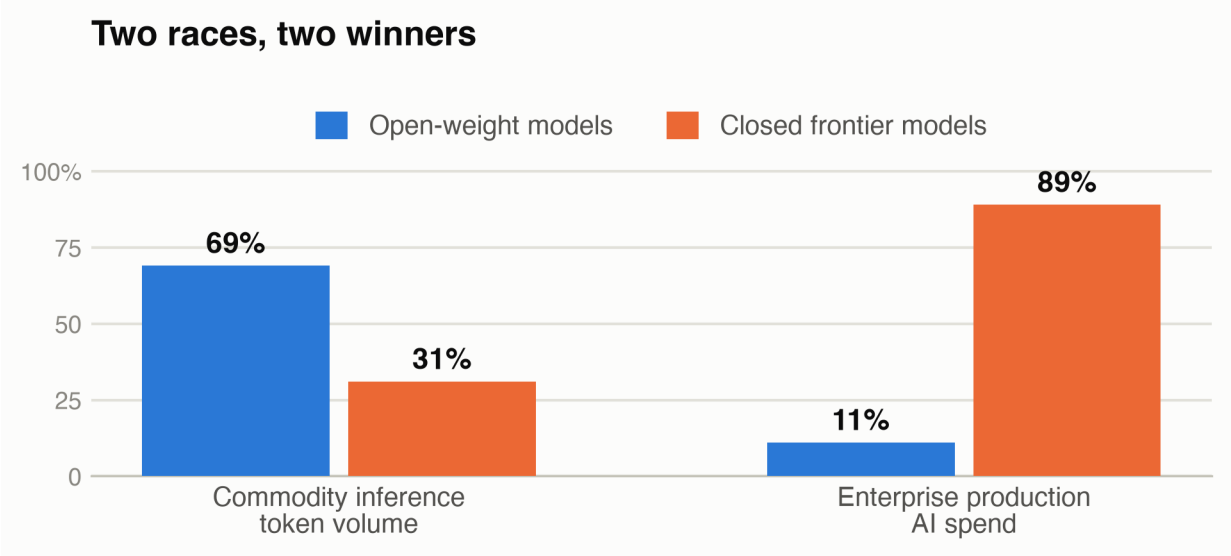


Figure 1: Two races, two winners

Why the moat moved off the model

The gap between open and closed frontier capability has held at roughly **four months** for a year running (three months in late 2025) – a stable, non-closing lead. Model weights are now a **perishable commodity**, not a durable asset: a four-month lead is worth nothing on a routine task and a large multiple on one where the model is the bottleneck for an expert's day – long-horizon agentic work, professional judgment, long-context retrieval. That is precisely where enterprises pay, and it explains the chart above: open models dominate usage, closed models dominate the P&L.

What frontier labs now race to acquire is not better weights but three inputs no release commoditizes: **proprietary domain data** (the operating history of a real business); **rigorous, task-specific evaluation** (what "good enough" means for a given job, not a leaderboard score); and **graded agentic trajectories** (verified, step-by-step records of real work, the training signal for the next generation of agents). The clearest proof of scarcity: a market worth an estimated $6B+ a year has emerged of vendors selling frontier labs synthetic, sandboxed replicas of enterprise software – fake Jira boards, fake EMRs, fake core-banking systems – purely to manufacture trajectories for post-training. Labs are paying real money to approximate what operating companies already possess natively. That is the moat, and it sits on operating companies' own books, not in a lab's weights.

Ecosystem leadership migrated geopolitically alongside this shift. Meta, once the undisputed open-weight leader, shipped a closed flagship model in 2026 and ceded that ground; Chinese labs now account for a plurality of open-model downloads, up sharply from a year earlier. Google, by contrast, leaned into permissive open licensing for its small/edge models – ceding the commodity tier on purpose to cement its position at the frontier.

What this changes in how we underwrite

Dimension	Thin AI wrapper	Data / evaluation moat
Owns	Rented intelligence + UI on a frontier API	Proprietary data, eval harness, graded trajectories
Defensibility	Erodes in ~4 months as rivals route elsewhere	Compounds – not reproducible by a competitor
Margin path	Exposed to vendor pricing swings	Improves via eval-driven routing to cheaper models
Our stance	Discount terminal value	Premium multiple for a durable, ownable asset

We no longer credit a target for "using AI," nor for which model it wraps – that layer commoditizes on a predictable clock. Three changes follow: we **discount terminal value** on wrapper-only revenue; we underwrite **model-routing savings** (~5x cost reduction on commodity tasks routed to open models) as realizable EBITDA upside, not blue-sky; and **no deal reaches IC** without an evidenced answer to *what does this company own that a frontier-lab release can't erase overnight?*

Diligence: six tests for any AI-exposed target

Test	What we require
1. Evaluation harness	A real eval suite run against the company's own workload, cost-per-task attached – not a vendor benchmark in a deck
2. Data provenance	Owned or exclusively licensed, not merely scraped or touched
3. Rented vs. defensible	Share of workflows on a single API; portability to a second model family without a rewrite

Test	What we require
4. Trajectory capture	Workflow logs retained, structured, and graded against verified outcomes – not just stored
5. Lab-licensing claims	Carried at zero until a signed, disclosed contract exists – no proven resale market yet
6. Log governance	Trajectory/prompt logs diligenced like IP: ownership, retention, DPAs, sub-processor terms

This is now a standing workstream in our diligence process on every deal with material AI exposure, with named owners and evidence requirements – not a bolt-on to IT diligence.

Where we hunt, and what we avoid

Commoditized model capability sharpens our existing lower-middle-market conviction. We favor **incumbent, "boring" operators** – specialty manufacturers, services firms, healthcare operators, logistics companies – sitting on years of proprietary operating data over AI-native startups pitched on branding; this is exactly the data frontier labs spend billions trying to synthetically approximate because they cannot buy or scrape the real thing. **Red flags:** single-vendor model dependency with no proprietary data; no internal evaluation capability; a commoditizable service wrapped in a chatbot; a thesis pitched on branding over dataset depth.

A second front: the compute layer

A parallel dislocation is forming one layer down, in the infrastructure that serves the models. Investor commentary we are tracking – an October 2026 X exchange between Chamath Palihapitiya and Elon Musk – describes hyperscalers as slow, bureaucratic compute buyers (months-long NDAs, multi-year escrow) losing "energized land" deals to faster operators, and warns that 2027 compute coming online faster than it can be absorbed will push hyperscalers to prioritize their own workloads and ration third parties through proprietary routers that take a margin – citing a rumored instance in Ant's IPO filing of a hyperscaler telling Ant it had no spare capacity. Musk adds a financing angle: debt markets have reportedly closed to AI-infrastructure borrowers without revenue independent of AI data centers – citing SpaceX's $40B GPU-related debt, backed by Starlink and launch revenue rather than the GPUs – and predicts consolidation among the 200+ neocloud operators that lack it.

(Source: Palihapitiya [@chamath] & Musk [@elonmusk], X thread, Oct 2026.) We treat this as directionally consistent with the thesis above – commoditized capacity rewards whoever holds a durable asset, data at the application layer, power/land/balance-sheet quality at the infrastructure layer – but flag it as one unverified social-media exchange between two investors with direct financial exposure (Chamath's own land portfolio benefits if this plays out; Musk's company is the credit example cited). We underwrite the pattern, not the 2027 date or the unnamed exception. **What this means for us:** for any target with AI-infrastructure exposure, we now diligence counterparty concentration and whether its AI-infrastructure debt is backed by revenue independent of the buildout itself.

Governance: the moat is also a liability filter

Open-weight proliferation brought a wave of derivative models with safety guardrails stripped out, circulating freely at scale. Enterprise capital has priced this in – governance maturity is now a buying criterion behind spend concentrating in better-governed closed labs. We treat this as a diligence lens: we document model provenance

and misuse exposure for every AI-enabled portfolio company; require consent, retention limits, IP treatment, and a named owner before any trajectory data is monetized; and underwrite without assuming the outcome of the live open-weight policy fight. We hold our own research to the same bar: no risk-mitigation credit for unverified practices.

Scorecard: what we believe, one year out

Claim	Status	Evidence
Open weights would lap closed models	Did not happen	Gap stable at ~4 months for a year
Nobody pays for closed models	Did not happen	Spend tripled to ~$37B; ~90% in three closed labs
Open source dominates enterprise	Reversed	Enterprise open-source usage fell 19% to 11%
Ecosystem ownership is the prize	Confirmed, changed hands	Meta ceded it going closed; China now leads
The model itself is the moat	False	Moat relocated to data, evaluation, trajectories

Supporting Analysis & FAQ

Q1: How reliable is the "~4 month" gap, and could it close tomorrow? It reflects a capability index tracked consistently for a year, widening slightly (three to four months) rather than narrowing – a steady-state window, not a trend to zero. Limitation: an index aggregates many benchmarks and can mask task-specific variance. Our response is not to bet on the index, but to require every target to run its own evaluation (Test #1).

Q2: Is the "data/evaluation moat" framing just a way to call AI wrappers overvalued? No – it is the mechanism. A wrapper's advantage is observably time-bound: once a cheaper model clears the bar, its edge evaporates within roughly the capability-gap window. Only the portion of value attributable to "which model we call" gets discounted toward zero – not owned data or graded operational exhaust.

Q3: Is trajectory data actually monetizable, or speculative? Demand evidence is solid – a multi-billion-dollar vendor market sells labs *synthetic* workflow approximations, which only makes sense if the real thing has real training value. We have no evidence of companies licensing *real* trajectory data under disclosed contracts, and flag that gap explicitly (Test #5). Our stance: capture internally; treat external licensing as unproven upside, not underwritten value.

Q4: Should we avoid thin-wrapper companies entirely? No – we discount, we don't exclude by category. A wrapper can migrate into the defensible category by building an eval harness and accumulating graded outcome data over time. We underwrite the moat that exists plus a credible path to build one.

Q5: What changes operationally, beyond the memo? Three things, already in effect: AI Capability & Data Assets is a standing diligence workstream with named owners on every AI-exposed deal; models carry a terminal-value discount on wrapper revenue and a margin line for routing savings; and sourcing now favors data-rich incumbents, redirecting Helios's effort.

Generated by SmallPE Agents — www.smallpe.com

Sponsored by inloop.studio